



# Technical Stack and Architecture for Real-Time Star Visualization App

## Primary Language and Rendering Stack

**Choice:** Use **Rust** as the primary language, with the **wgpu** graphics API for rendering (via `wgpu` crate) and compute shaders. This stack best meets our cross-platform, performance, and maintainability needs:

- **Cross-Platform & Performance:** Rust's `wgpu` is a safe, modern wrapper over Vulkan/Metal/DirectX, running natively on all major OSes <sup>1</sup>. Applications using `wgpu` automatically work on Vulkan (Linux), DirectX 12 (Windows), Metal (macOS), etc., giving first-class support across platforms <sup>2</sup>. This ensures Windows, macOS, and Linux are all supported with one codebase. `wgpu` also supports WebGPU/WebGL for web, though our focus is desktop. Crucially, `wgpu` supports GPU **compute** and graphics, suitable for both rendering and parallel simulations on the GPU <sup>3</sup>.
- **Performance and Safety:** Rust delivers C/C++-level performance while preventing memory bugs, crucial for a complex high-fidelity app <sup>4</sup> <sup>5</sup>. No garbage collector or heavy runtime – Rust's zero-cost abstractions and compile-time checks mean we can push the GPU hard (60 FPS at 4K) without sacrificing reliability. We avoid crashes and memory leaks that could occur in C++ (especially under multi-threading/GPU stress) thanks to Rust's safety guarantees <sup>6</sup>.
- **Maintainability:** Using `wgpu` (which is based on the WebGPU standard) keeps the code simpler than raw Vulkan. The WebGPU API is designed to be a cleaner, modern graphics API that requires less boilerplate than Vulkan <sup>7</sup>. This means faster development and easier long-term maintenance. We get most of Vulkan's performance without managing every low-level detail ourselves. In contrast, a C+++Vulkan solution would demand significantly more code (and debug time) for little gain. Similarly, **Godot 4** or **Bevy** were considered, but a custom Rust/wgpu stack gives us finer control over a specialized rendering pipeline. Godot's engine imposes constraints and extra overhead (and using GDScript/C# would hurt performance). Bevy (Rust) uses wgpu under the hood, but introducing a full ECS engine would add complexity and potentially impede the custom multi-pass, multi-layer rendering we need. A lean custom engine built directly on `wgpu` keeps our distribution simple (just a single Rust binary + assets) and our control maximal.

**Conclusion:** **Rust + wgpu** is the decisive choice for the core. It provides a portable, high-performance graphics and compute platform in one package <sup>3</sup>, aligning with our offline, cross-platform requirement. We will use the Rust package manager **Cargo** for build and dependency management (MIT/Apache-2 licensed ecosystem). The initial project setup should be a new binary Rust crate:

```
cargo new star_viz --bin
```

Then add these core dependencies in `Cargo.toml`:

```
[dependencies]
wgpu = "0.27"          # GPU rendering and compute (WebGPU native API) 1
winit = "0.30"          # Window creation and input handling 8
```

We also include `pollster = "0.3"` (to easily block on async GPU init), and logging crates (`log = "0.4"`, `env_logger = "0.9"`) for debug output. With this foundation, we have a robust, cross-platform engine ready for real-time graphics and computation.

## Rendering Architecture

We will implement a **custom multi-pass rendering pipeline** in wgpu to simulate the Sun's layers and phenomena, integrating compute shaders for physics-heavy updates. The rendering is organized into distinct passes for different components, all rendered in **HDR (high dynamic range)** and then tone-mapped for display:

- **Photosphere Pass:** Render the star's surface (photosphere) as a sphere mesh with a high-detail procedural texture. We'll use a shader that generates procedural **granulation** (turbulent cell-like pattern) using noise algorithms. Instead of simulating convection, we tune procedural noise parameters (cell size, contrast) to match observed granulation statistics (approx. 1000 km cells, ephemeral patterns) – providing a realistic visual **approximation** of the boiling surface without heavy computation. This pass outputs the base color and intensity of the Sun's disc.
- **Chromosphere & Corona Passes:** Above the surface, we render atmospheric layers:
  - The **chromosphere** (thin inner atmosphere) can be drawn as a shell around the sphere with additive blending – a semi-transparent glow, possibly using a procedural texture (for spicules or mottling) tinted reddish.
  - The **corona** (outer atmosphere) is rendered using volumetric techniques. We employ a **ray-marching shader** in a fullscreen pass: for each view ray through space, sample an analytical density function for plasma as a function of altitude and latitude. We use an **empirical density/temperature profile** (e.g. an exponentially decreasing density with scale height ~50,000 km) rather than solving fluid equations. This produces a halo around the Sun that brightens near the limb. Because coronal brightness spans orders of magnitude, we keep this in HDR and let tone-mapping handle the extreme dynamic range. No full MHD simulation is needed – we plug in known density formulas to get a visually convincing corona in real-time (philosophy: use known physics *results* directly rather than compute them from scratch each frame).
- **Magnetic Field Lines (PFSS) Pass:** To visualize the Sun's magnetic field, we draw **field lines** as tubular lines or curves. We adopt the **Potential-Field Source-Surface (PFSS)** model – a widely used, computationally inexpensive model that captures the large-scale coronal field structure <sup>9</sup>. PFSS assumes the corona is current-free (valid most of the time, giving a reasonable field estimate <sup>10</sup>). We will load or generate a low-resolution solar magnetogram (surface field map, possibly an included sample or a simple dipole for presets) and compute field lines from it:

- We launch a **compute shader** that seeds many field lines (e.g. 10,000) at random or user-selected footpoints on the photosphere. Each thread traces a line by iteratively stepping along  $\mathbf{B}$  (the magnetic field vector), using the PFSS field equation (solving  $\nabla^2\Psi=0$  via spherical harmonics or a precomputed grid) <sup>10</sup>. Because PFSS solutions can be computed quickly (even via spherical harmonic coefficients), we can either precompute the field on CPU at load or evaluate in the compute shader on the fly. The result is a set of polylines (lists of 3D vertices) for magnetic loops.
- The field lines are then rendered in a dedicated pass. We'll represent them either as line primitives or thin triangle strips (for better lighting). A geometry shader is not available in wgpu/WGSL, so for better visuals we may generate cylinder or tube meshes for lines if needed. Lines will be colored (e.g. closed loops in white/gold, open field lines in different colors for polarity) and added with additive blending or bright emissive shading so they glow on top of the corona.
- *LOD Strategy:* To maintain performance, we can adjust the number of field lines based on view. For a full-disk view or heliospheric view, we'll render fewer, broader-strokes field lines (e.g. only the largest loops or a random subset of the 10k) to avoid overdrawing 10k lines in the distance. When zoomed in to an active region, we can render a denser set of local field lines for detail. This Level-of-Detail scaling ensures we meet 60 FPS while always showing appropriate detail.
- **Particle and Event Effects Pass:** Dynamic phenomena like solar flares, prominences, and coronal mass ejections (CMEs) are handled by GPU **particle systems**:
  - We allocate particle pools for things like eruptive plasma or glittering flare ribbons. Their motion is updated with compute shaders each frame: e.g. integrating simple equations or following field lines. For example, a flare might emit particles that travel along field lines (computed from the PFSS field) to show hot plasma motions. CMEs can be a burst of particles expanding outward.
  - These particles are rendered as billboard sprites (textured quads always facing camera) or point sprites, with additive blending to represent glowing plasma. The particle update compute shader can also apply simple physics (gravity, drag) or keyframe animations. This approach offloads thousands of particle updates per frame to the GPU for efficiency.
  - We will **pre-script** certain events rather than fully simulate physics: e.g. when the user triggers a flare, we initiate an animation where brightness and particle emission follow a time curve from real observations (rise and decay times, etc.), rather than simulating magnetic reconnection from first principles. This yields a visually authentic animation with minimal computation.
- **UI & Overlay Pass:** Finally, we render the UI on top (discussed in next section) and any on-screen text/labels. The UI will render in its own pass using its library's drawing routines (e.g. ImGui or egui draws to an offscreen texture or directly to the swapchain as an overlay). We ensure the UI is drawn *after* tone-mapping (so that the UI is crisp and not affected by the HDR process).
- **HDR and Tone Mapping:** All 3D scene rendering passes (photosphere, corona, field lines, particles) will render into an **HDR framebuffer** (e.g. `wgpu::TextureFormat::Rgba16Float`) so we can handle the Sun's extreme brightness range without clamping <sup>11</sup>. After composing the scene in HDR, we perform a full-screen post-process pass with a tone-mapping shader. We'll use the ACES filmic tone mapping curve (a standard used in games/film for realistic highlights roll-off <sup>12</sup>) to convert the HDR image down to SDR (sRGB) for display. This ensures features like the bright corona and flares can be depicted with bloom and intensity but still viewed on a standard monitor. We also enable

**bloom/glow** as needed by blurring bright parts of the HDR image and adding them (this can be an additional post-process pass if time allows, to enhance the brilliance of flares and corona).

- **Depth Buffer & Blending:** We'll use a depth buffer for the 3D geometry (sphere, lines, particles) so that, for example, field lines properly appear emerging from the surface and going behind the sun as needed. Some passes (like corona volumetric) might not write depth (they are more like an atmosphere overlay). We will carefully order and blend passes:

- Opaque pass (photosphere) with depth test.
- Corona volumetric (additive blending, no depth write, but depth test to fade out behind sun).
- Field lines and particles (additive/emissive blending, depth test on, so they can go behind the sun or each other appropriately).
- UI (always on top in screen space).

This multi-pass setup, combined with compute shaders for intensive tasks, will achieve a realistic yet performant visualization. By **baking in known physics approximations (PFSS field, empirical density profiles, etc.)**, we avoid heavy real-time simulations and keep frame rendering within budget. The PFSS model in particular provides a "good enough" magnetic field visualization at a fraction of the cost of full MHD <sup>10</sup>, aligning with our goal of **visual accuracy over physical accuracy** for outreach and education.

## UI Framework and Panels

For the UI, we choose an **immediate-mode GUI** library that integrates with Rust and wgpu, providing native **dockable panels** and custom widgets. The recommended choice is **Dear ImGui** via Rust bindings (`imgui-rs`), with docking enabled, rendered using wgpu. ImGui offers a polished, high-performance UI with built-in support for docking windows/panels – exactly what we need for a professional, resizable panel interface. An alternative was **egui**, a pure-Rust GUI; however, egui lacks built-in docking (it requires an add-on <sup>13</sup>), whereas ImGui's docking is battle-tested. We opt for ImGui for a snappier dev experience and its flexible panel system.

**Integration:** We add the following UI crates to Cargo dependencies:

```
imgui = { version = "0.10", features = ["docking"] }           # Dear ImGui Rust
bindings (enable docking branch)
imgui-wgpu = "0.16"                                         # Renderer to draw
ImGui via wgpu 14
imgui-winit-support = "0.10"                                    # Integrates ImGui
with winit (event handling)
```

These crates let us create an ImGui context, build UI panels, and render them as part of our wgpu frame. ImGui's docking system will allow our app to have multiple panels (as tabs or floating windows) that the user can rearrange and resize <sup>13</sup> – perfect for a complex visualization tool.

**UI Elements:** We will design a **dockable panel layout** such that users can organize controls on either side of the 3D view or even detach them: - **Wavelength / View Mode Toggles:** A panel (or toolbar) with checkboxes or radio buttons to switch the visualization mode. For example, toggle different wavelengths

( $171\text{\AA}$ ,  $304\text{\AA}$ , etc. if we simulate different filters), or toggle viewing the magnetogram, corona, field lines on/off. These toggles will call into the app state to enable/disable rendering layers. - **Star Preset Selector:** A dropdown or list to select different star or Sun presets (e.g. "Quiet Sun", "Active Sun with big sunspots", "Other star (with different parameters)"). This triggers loading a preset file (see Data Management). - **Time Controls:** Buttons or slider for time animation – e.g. play/pause rotation or advance simulation time. Since the app can animate (rotation of the sun, loop oscillations, etc.), a time slider allows scrubbing back and forth. We will implement a simple time accumulator in the app and use UI controls to modify it or pause it. - **Camera Controls:** Although the primary camera control will be mouse-driven (click and drag to rotate around the sun, scroll to zoom, etc. via winit events), we will also provide UI sliders for camera latitude/longitude and distance, for precise adjustments. A "reset camera" button also goes here. - **Parameter Sliders:** Many visual parameters (and a few physics knobs) will be adjustable with sliders. For example: granulation contrast, corona brightness scale, number of field lines shown, particle density, flare speed, etc. Each slider directly binds to a variable in our program state. ImGui makes it easy to create sliders for floats/integers and we'll group them in collapsible sections (e.g. "Corona settings", "Field lines settings", etc.). This invites user exploration of the visualization. - **FPS/Performance Display:** We can include a small overlay (perhaps in a corner) showing current FPS and memory usage, to help tweak performance. ImGui can display text or graphs easily.

The UI will follow a **dockable panel architecture** – for instance, the left side can have a tabbed panel (one tab for "Layers", one for "Camera", etc.), and the bottom could host a timeline panel, etc. Users can drag these out into separate windows if they have multiple monitors or wish to rearrange <sup>13</sup>. The immediate-mode nature of ImGui means we will construct the UI every frame in code, reflecting current state.

**Maintainability:** ImGui's approach lets us quickly add new controls. We will organize UI code in its own module (e.g. `ui.rs`) which renders the panels given the current application state (passed as mutable reference). The separation ensures the rendering logic (in the core engine) is decoupled from UI logic.

Finally, because ImGui is rendered using our own wgpu render pass (via `imgui-wgpu`), it plays nicely with our HDR pipeline. We will render the UI *after* tone-mapping, in standard sRGB space, so the UI colors aren't affected by scene brightness. The result is a responsive, professional UI with flexible panel management, fulfilling the requirement for dockable/resizable control panels.

## Data and State Management

Our application will manage simulation configuration and state in a transparent, user-friendly way, using simple data formats (JSON) for presets and saves, and carefully dividing what runs on CPU vs GPU for efficiency.

- **Preset System:** We will support loading star preset files that define initial conditions and parameters (e.g. a preset for the Sun in a quiescent state, one for a flaring state, perhaps presets for different star types). We recommend using **JSON** for these preset files – it's human-editable and widely understood, and our Rust code can easily parse it via Serde. Each preset JSON can include:
  - Star parameters: radius, surface gravity (for plasma scale heights), approximate magnetic field map (we could include a small array or an ID of a bundled magnetogram image).
  - Visual settings: e.g. color of the chromosphere, brightness multipliers for corona, etc.
  - Predefined events: e.g. a preset might include a scripted flare at a certain time.

We will create a `Preset` Rust struct (with Serde `Deserialize`) to load these. Example format:

```
{  
    "name": "Active Sun with big sunspot",  
    "radius": 696340.0,  
    "magnetogram": "data/sunspot_region.fits",  
    "corona_brightness": 2.0,  
    "field_line_footpoints": [[-30.0, 45.0], [10.0, -20.0], ...]  
}
```

This way, Claude (the developer) or users can add new presets by editing JSON rather than recompiling code. We'll bundle a few JSON presets with the app for out-of-the-box use.

- **External Data (Optional FITS import):** While the app is **offline-first** (no network needed), advanced users might import real scientific data files they have on disk. We will provide an *optional* path to import **FITS** files, the standard astronomy format, for things like magnetograms or images. Using the Rust `fitsio` crate (optional dependency), we can read a FITS image (e.g. an HMI magnetogram of the Sun's surface) and use that as input for the field line computation or to texture the photosphere. This is a secondary feature – not required for core usage, but it extends the tool's utility for those who have data. We won't support writing out complex FITS or any live data APIs (no calls to VSO or Heliovewer) to keep things simple and offline. If a user wants latest data, they can download a FITS themselves and import it.
- **Simulation State (CPU vs GPU):** We delineate clearly what data is stored and updated on CPU vs GPU for performance:
  - **CPU State:** The CPU will hold high-level state and low-frequency data. For example, current time/frame counter, user-selected preset parameters, and any data that is cheap to update or needed for UI logic. The CPU will also manage any one-time computations like loading a preset, reading files, initializing the PFSS model (if done on CPU), etc. We'll use Rust structs to hold this, e.g. an `AppState` struct with fields for each subsystem (camera, layers toggles, etc.).
  - **GPU State:** The heavy lifting resides on the GPU. This includes large arrays or textures: e.g. the 3D volume or function for corona emission (if we precompute a 3D texture of emissivity, we'd store that on GPU), the vertex buffers for field lines and particle systems, and uniform buffers with parameters for shaders (matrices, time, scales). The **compute shaders** will update GPU buffers directly each frame (for particles), minimizing CPU-GPU data transfer. We will use wgpu's **storage buffers** for things like particle positions and field line vertices, so that compute shaders can write to them and render shaders can read from them. The CPU just orchestrates when to run those computes.
- **Update Frequencies:** Most per-frame updates (60 fps) happen on GPU: e.g. particle integration every frame. The CPU will tick at 60 fps as well but doing light work (polling input, updating the ImGui UI, issuing draw calls). Some things update at lower rates: e.g. recomputing all magnetic field lines from scratch may only be triggered on preset load or if the user toggles a new magnetogram (since 10k field lines integration is maybe a few milliseconds on GPU, we can do it occasionally but not every frame). We can also spread out updates: e.g. update 1/10th of the field lines per frame on a rolling basis if we want a continuously evolving field visualization without a big one-frame cost.

- **Save/Load Visualization State:** Besides presets (which are like templates), the user can save the **current state** of the visualization (including any interactive changes) to disk. This would also be a JSON file (or could reuse the preset format). Essentially, we serialize the `AppState` (which includes which preset was base, plus any tweaks, camera position, time, etc.) using Serde. This allows a user to set up a nice scene (e.g. at a particular time with a particular camera angle and event happening) and save it, then later reload exactly that scene. Since JSON is our choice, this also means saved states are human-readable and editable.
- **Screenshots and Export:** Meeting the secondary requirement, we'll implement a screenshot function to capture the current framebuffer to a PNG image. Using the `image` crate, we can take the `final rendered image` (after tone mapping) from the swap chain (via `CommandEncoder::read_texture` or similar) and save as PNG. The screenshot will include some metadata (maybe in filename or a small JSON sidecar) like the time and preset name. We may also allow exporting the current magnetic field lines data – for example, saving them as a simple CSV of XYZ points or Wavefront OBJ format, so that an external 3D tool or 3D printer software could import the magnetic loops. This is optional and straightforward to add using our existing data (just writing out the vertex lists we already have).

**No Live Data, Fully Offline:** Importantly, *no* part of the app requires internet. All data is either procedurally generated or loaded from local files. This guarantees the app works in classrooms, museums, or remote locations with no connectivity. We avoid dependency on external APIs that could break or limit usage. The app as distributed will include all necessary data in an `assets/` folder (e.g. sample magnetogram FITS, preset JSONs, maybe a couple of representative images). The user's GPU does all the work – for instance, computing PFSS loops for a provided magnetogram, or generating noise for granulation – ensuring the experience is self-contained.

**Data Formats:** In summary, we use **JSON** + Serde for config/presets and state save, **FITS** (via `fitsio`) optionally for astronomy image inputs, and possibly simple text formats for trivial exports. All of these are open formats aligning with our open-source nature. The emphasis is on *readability and simplicity* – e.g., JSON for states so that an educator could tweak a parameter by editing a number in a file if needed.

By structuring data and state this way, we ensure the simulation runs smoothly at 60fps with minimal stalling: heavy calculations are either done upfront (load time) or steadily on the GPU, while the CPU remains free to handle user input and UI. The end result is a responsive application where users can seamlessly load different scenarios, adjust parameters, and even integrate their own data, all while offline.

## Project Structure and Module Organization

We will organize the project for clarity, modularity, and ease of contribution. Using Cargo's conventions, we create a clean source layout:

```
star_viz/
├── Cargo.toml
└── src/
    └── main.rs      # Program entry, initializes window, wgpu, and enters
                      main loop
```

```

├── renderer.rs      # Rendering pipeline setup and per-frame rendering code
├── shaders/         # Directory for shader source files (WGSL)
|   ├── photosphere.wgsl
|   ├── corona.wgsl
|   ├── fieldlines.wgsl
|   ├── particle_update.wgsl
|   └── tonemap.wgsl
├── simulation.rs    # Physics approximations, e.g. PFSS field model, data
for density profiles
|   ├── ui.rs          # UI layout and event handling (Dear ImGui integration
and panel definitions)
|   ├── state.rs        # Definition of AppState, Preset struct, and
serialization (Serde)
|   └── data/           # (Optional) Embedded data files or lookup tables
(maybe included via build)

```

Explanation of key modules:

- **main.rs:** Sets up the application window (using winit), initializes the wgpu **Device**, **SwapChain**, etc., and configures ImGui. It will load an initial preset (e.g. default Sun) and then enter the main event loop. In the loop, we handle input events (keyboard, mouse via winit), update the **AppState**, run any scheduled compute tasks (like particle updates), then call into **renderer.rs** to draw the frame, and finally render the UI via **ui.rs**. Main.rs coordinates these pieces.
- **renderer.rs:** Responsible for graphics init and drawing. It will create all the **wgpu pipeline objects**: the render pipelines for each pass, bind group layouts and bind groups (for uniforms, textures), and the depth texture. This module will likely define a **Renderer** struct holding things like the **wgpu::Device**, **Queue**, pipeline states, and resource handles (textures, buffers). It will also contain the **render\_frame()** function that encodes all rendering passes each frame in the correct order. Rendering code is here, separate from simulation logic, to keep concerns separate. This is where we load our WGSL shaders (from files in **src/shaders/**) and create pipeline layouts. By isolating shader code in **.wgsl** files, we make editing and experimenting with shaders easier (no need to recompile the whole program to tweak a shader; they can be reloaded if we implement hot-reloading for development).
- **shaders/:** Contains WGSL shader source files for each stage of our pipeline (vertex/fragment for the photosphere, compute for particles, etc.). Organizing them in a folder makes it clear and allows syntax highlighting in editors. These will be included in the binary with **include\_str!** or similar at compile time, or we can load them at runtime for easier iteration. In any case, keeping them separate keeps Rust code cleaner.
- **simulation.rs:** This module implements the **PFSS field computation** and other physics helpers. For example, it might have a function **generate\_field\_lines(magnetogram\_data) -> Vec<LineVertex>** that either runs on CPU or dispatches a GPU compute and then reads back results. It can also include any code for procedural generation (noise functions for granulation if we use CPU for that, or generation of initial particle distributions, etc.). Essentially, anything that is

about preparing or computing the data that gets rendered (but not the rendering itself) lives here. This might be further split into submodules if it grows (e.g. `field.rs` for magnetic field math, `plasma.rs` for density functions). It will likely use math libraries (like `glam` for vector math) and possibly `rayon` if we do CPU parallel loops.

- **ui.rs:** Contains all Dear ImGui calls to construct the interface. We will define functions for each panel, e.g. `ui::show_controls_panel(&mut imgui::Ui, state: &mut AppState)`, `ui::show_info_panel(...)`, etc. This module will handle mapping user interactions to changes in our `AppState`. For example, if the user toggles a checkbox for "show corona", the UI code will set `state.render_corona = true/false`, which the renderer will read and act on. Keeping UI separate means if in the future we switch UI libraries or need a different layout, we can change this module without touching core logic.
- **state.rs:** Defines the central `AppState` struct that holds all runtime state (e.g. booleans for which layers are on, current simulation time, current camera parameters, etc.), as well as structures for **Presets**. It will implement Serde `Deserialize/Serialize` for those, so that loading a preset JSON into a `Preset` struct is easy, and saving state to JSON is just serializing `AppState`. This module may also contain the logic to apply a Preset to the running state (e.g. a function `apply_preset(state: &mut AppState, preset: Preset)` that sets up all parameters and possibly triggers recomputation of field lines). We isolate this to keep configuration management in one place.
- **data/:** Not a code module, but a place to store static data files if any. For instance, if we include a default magnetogram image or a lookup table (like a precomputed 1D density profile table), we can put it here and use `include_bytes!` or have the build script bundle it. This keeps the project structured (and can be packaged with the binary on install).

**Build System:** We use **Cargo** exclusively. The project will be a single Cargo package (no need for a workspace unless we later split into library crate vs GUI crate, etc., but initially one crate is fine). The build will produce one binary (e.g. `star_viz.exe` on Windows). We will enable relevant features in Cargo for dependencies (e.g. for `wgpu` we may enable the "spirv-cross" or shader translation features if needed, or for `imgui` the "docking" feature as shown). If we decide to allow the optional FITS support, we can include it as `[dependencies.fitsio] optional = true` and use Cargo features to toggle it. But since it's a small crate, we might just include it by default for simplicity.

**Dependency Management:** All dependencies are from crates.io (open source). We have already listed the major ones: `wgpu`, `winit`, `imgui` + integrations, `serde` + `serde_json`, `glam`. We'll pin versions known to be stable. We should also add:

```
serde = { version = "1.0", features = ["derive"] }
serde_json = "1.0"
glam = "0.30"
fitsio = "0.21"          # (optional, for FITS import)
image = "0.24"           # (for PNG screenshot encoding)
```

And possibly `rayon = "1.7"` if we plan to use it for any parallel CPU computations (like computing thousands of field lines in parallel on CPU if we ever do that instead of GPU). However, most heavy lifting is on GPU, so `rayon` is optional.

This structured approach ensures that **Claude (the developer)** can easily navigate the codebase: rendering code vs simulation code vs UI code are cleanly separated. It will also facilitate future contributors (the separation of concerns is clear). The directory layout and module names can be adjusted as needed, but the key is to avoid one monolithic file – instead we want logically divided components with well-defined interfaces (e.g. the `Renderer` struct might expose `fn render_frame(&mut self, state: &AppState)` and the UI might expose `fn ui_frame(ui: &mut imgui::Ui, state: &mut AppState)` etc.).

We will use **Cargo features** to conditionally compile certain things if needed (for example, a “dev” feature might enable an debug camera or extra UI, or an “import-fits” feature to include FITS support to trim release builds). Initially, though, simplicity wins: get the core working, then iterate.

Finally, regarding the **build and distribution**: We target a static or nearly-static binary. Rust’s cargo will produce a single binary that we can ship. Any asset files (like shader WGSL or preset JSON) can be embedded (via `include_str!` / `include_bytes!`) to truly make it a single file application. Alternatively, we ship an `assets/` folder alongside if we prefer easier modding (users could then edit the JSON or shaders). Either approach satisfies offline use – no further installation or internet needed once the user has the binary + assets.

## Performance Targets and Optimization Strategy

We design the application from the ground up to hit **60 FPS at 4K resolution** with all features (full corona rendering, ~10k magnetic field lines, particle effects) on a modern mid-range GPU (e.g. NVIDIA RTX 3060). Our strategy to achieve this involves efficient use of the GPU, budgeting of GPU memory, and a robust profiling regimen during development:

- **GPU Budgeting:** We will manage GPU memory to stay within a comfortable budget (targeting <4 GB VRAM usage even at 4K). Key allocations:
- **Render targets:** The HDR framebuffer at 4K (RGBA16F) will consume ~4K4K8 bytes ≈ 128 MB, which is fine. The depth buffer (4K, 32-bit) ~ 64 MB. These are significant but only two.
- **Textures:** We might have a few textures (e.g. a 4K texture for photospheric detail or a sunspot map). These will be at most tens of MB. We avoid gigantic textures – procedural generation covers most detail, so we don’t ship dozens of large images.
- **Buffers:** Field line vertex buffers: suppose 10k lines \* 100 vertices each = 1,000,000 vertices. If each vertex is position (12 bytes) + color (4 bytes), ~16 bytes, that’s ~16 MB – trivial for modern GPUs. Even if doubled for two ends or normals, still < 32 MB. Particle buffers: if we have say 20k particles, each with position, velocity, etc., maybe 64 bytes each, that’s ~1.2 MB. Negligible. Uniform buffers and small stuff are in the KB range. So geometry and particles are easily within budget.
- **Coronal volume:** If we did decide on a 3D volume texture for corona (say  $256^3$  grid of density values), that’s 16M cells; at 1 float each (~64 MB). We can also compress or procedurally generate on the fly. But more likely, we avoid a full 3D texture and compute density in shader procedurally (which saves memory entirely). So we keep an eye on such potential big allocations. Overall, a few hundred

MB total usage is expected, well under an RTX 3060's 12 GB. This leaves headroom for when running on lower cards (e.g. 4 GB cards should still manage if resolution or certain details are lowered).

- **Maintaining 60 FPS:** We leverage the GPU heavily to ensure the CPU is not a bottleneck. The CPU will just record commands and feed data; it won't, for example, loop over 10k field lines itself per frame – that's done in a compute shader in parallel. Also, we design all our passes to be parallel-friendly on GPU:
  - The field line rendering and particle rendering are largely **fill-rate bound** (drawing lots of small primitives). We will measure and if needed, adjust. For example, if  $10\text{k field lines} \times 100 \text{ pts} = 1 \text{ million vertices}$  is too slow to rasterize at 4K (though likely fine for a 3060), we could reduce lines when far as discussed, or use simpler representations. We can also experiment with drawing field lines as thinner lines (less pixel coverage) or even as impostors.
  - The corona ray-march shader could be expensive at 4K (each pixel does many samples). To keep 60fps, we will optimize the shader: use a coarse step size that still looks good, possibly ray-march only a certain distance out. We could also implement **adaptive sampling**: e.g. fewer samples in darker regions. If needed, we can dynamically reduce corona rendering resolution at runtime (render the corona to a half-res texture and upscale with a blur) – a common trick for expensive fullscreen effects. But we will first try to meet performance at full res.
  - We will use **Level of Detail (LOD)** dynamically: As noted, reduce complexity for far zoom (skip granular details, reduce particle counts when the whole Sun is in view, etc.). These can be tied to camera distance or a user setting for quality vs performance.
  - The UI (ImGui) is lightweight (a few hundred draw calls at most) and shouldn't impact 60fps significantly. We will, however, be careful to only redraw what's needed; ImGui by default redraws every frame (immediate mode), which is fine given its small cost relative to the 3D rendering.
- **Profiling and Optimization:** Claude should build with a mindset of measuring performance early and often. We will integrate basic profiling tools:
  - Use wgpu's built-in support for capture and debugging. We can enable `wgpu::Device` features to capture GPU timings or insert debug markers around passes. Also, using tools like **RenderDoc** or NVIDIA Nsight, we will capture frames to see where GPU time goes (e.g. how long the corona shader takes, or if the field line drawing is a bottleneck) – these external GPU profilers are essential for pinpointing bottlenecks in a frame <sup>15</sup>.
  - For CPU profiling, we can use lightweight crates like `tracing` or `puffin` to instrument our code. For example, we can mark the UI construction, physics update, etc., and ensure none of those are taking too long. Ideally, the CPU frame time should stay low (a few ms) so the GPU is the limiting factor (GPU likely taking ~16ms at 60fps).
  - We'll monitor memory usage and frame times using the UI itself (we can display them). During development, running in debug mode with smaller window and then in release mode at 4K will help test the performance envelope. We should also test on multiple GPUs if possible (NVIDIA, AMD, and an Intel iGPU) to ensure we don't accidentally use something non-portable or super heavy on certain drivers.
  - **Optimization techniques:** If a particular pass is too slow, we'll optimize shader code (e.g. reduce divisions, use approximations, precompute lookup tables). For instance, if computing `exp()` for corona density is slow, we can approximate with a cheaper function or a small 1D texture. We will

also take advantage of WGSL's capabilities: e.g. use 16-bit float math if possible for some calculations to speed them up, use shared memory in compute shaders for efficient data sharing if needed, and avoid unnecessary memory transfers (keep data on GPU once generated).

- We'll keep frame latency low by using *wgpu's default of triple buffering*. We also consider using multiple command buffers (one for each pass) recorded in parallel threads if wgpu supports it (wgpu does allow multi-threaded recording via multiple encoders). But unless needed, a single-threaded command encoding likely suffices at 60fps.
- **Meeting the 60fps@4K Goal:** Through the above measures, we are confident the target is achievable. Similar graphics (space sims, solar visuals) have been done on lesser hardware. Our content is largely procedural and vector-based (lines), which is easier than, say, hyper-realistic 3D models. The key is to keep an eye on the heaviest operations (the volumetric shader and massive overdraw of lines/particles). With dynamic LOD and fine-tuning, we'll ensure those stay within budget. For example, if a user zooms all the way out to see the heliosphere, we might disable some minor effects (granulation won't be visible anyway) and limit particle counts to maintain FPS.
- **Testing and Tuning:** We will include a **benchmark mode** in development that can stress-test the scene: e.g. render maximum 10k lines, maximum particles, and print the frame time. This helps determine if any subsystem needs further optimization. Using conditional compilation, we can turn off V-sync and push frames as fast as possible to measure the theoretical max throughput.

In summary, by leveraging Rust+wgpu's performance, doing heavy computations on GPU, and adopting adaptive detail techniques, we will hit smooth 60fps visuals. Memory won't be an issue given our mindful budgeting, and any emerging bottlenecks will be addressed with the aid of profiling tools and adjustments. The end result will be an application that feels fluid and responsive even with all the eye-candy enabled, on a typical gaming PC – fulfilling the goal of **stunning real-time visualization** with no compromises on frame rate.

## First Steps for Claude (Developer Instructions)

To kick off development with the chosen tech stack, Claude should perform the following:

1. **Create the Project:** Use Cargo to create a new binary project:

```
cargo new star_viz --bin  
cd star_viz
```

This will create the basic Cargo.toml and src/main.rs.

2. **Add Dependencies:** Open `Cargo.toml` and add the decided dependencies under `[dependencies]`:

```
[dependencies]  
# Windowing and Graphics
```

```

winit = "0.30"
wgpu = "0.27"
pollster = "0.3"
# UI Libraries
imgui = { version = "0.10", features = ["docking"] }
imgui-wgpu = "0.16"
imgui-winit-support = "0.10"
# Data formats and math
serde = { version = "1.0", features = ["derive"] }
serde_json = "1.0"
glam = "0.30"
# Image and FITS (optional)
image = "0.24"
fitsio = "0.21"
log = "0.4"
env_logger = "0.9"

```

Then run `cargo check` to fetch and ensure all dependencies compile. This set gives us:

3. winit for window+input, wgpu for GPU, pollster to simplify async,
4. Dear ImGui with docking (via imgui crate) and the wgpu + winit integration,
5. Serde and serde\_json for config, glam for linear algebra,
6. image for screenshot saving, fitsio for optional FITS import,
7. log/env\_logger for logging debug info.
8. **Project Structure Setup:** Create the module files as planned. For now, Claude can focus on stubbing out the major modules:
9. In `src/main.rs`: create a basic window using winit and initialize wgpu (adapter, device, queue, swapchain). Set up env\_logger for debugging. Also initialize ImGui context and the imgui-wgpu renderer.
10. Create `src/renderer.rs`: define a `Renderer` struct with placeholders for pipelines, and perhaps a `init(device: &wgpu::Device, config: &wgpu::SurfaceConfiguration) -> Renderer` and a `render_frame(&mut self, state: &AppState, imgui_draw_data: Option<&imgui::DrawData>)` that will be fleshed out later. For now it can clear the screen to a color and draw nothing.
11. Create `src/ui.rs`: define a function to configure the dockspace (ImGui docking root) and a placeholder panel (e.g. show "Hello, Sun!"). Get the docking UI set up early to verify ImGui with docking works (there are examples in imgui-rs for enabling docking, which usually involves calling `io.config_flags |= DockingEnable`).
12. Create `src/state.rs`: define an `AppState` struct with a couple of example fields (like `pub show_corona: bool` etc.) and derive Serde. This will expand as we integrate.
13. Other files (simulation.rs, shaders/) can be added a bit later once the basic window+triangle is running.

14. **Hello Triangle Test:** As a quick graphics sanity check, have the renderer draw a single triangle to ensure wgpu is set up correctly. Then integrate ImGui overlay on top (maybe show the FPS in ImGui). Getting this minimal frame running verifies the scaffolding. Use `wgpu::SurfaceConfiguration` with an SRGB format for now (we will later switch to HDR buffer for tone mapping once we implement that).

15. **Implement Incrementally:** Once the basic loop is running:

16. Implement the photosphere rendering (create sphere mesh, a simple shader).
17. Then add corona effect, field lines, etc., one by one, profiling as needed.
18. Use the UI to toggle these on/off to test their performance impact early.

By following this plan, we've chosen a powerful and focused tech stack and provided Claude with concrete starting instructions. The decisions (Rust/wgpu, custom pipeline, Dear ImGui UI, JSON presets, offline data) are all aligned with the project constraints and goals, ensuring development proceeds decisively with the right tools for the job. This technical specification should guide Claude in building a high-performance real-time star visualization application with confidence and clarity.

---

1 2 GitHub - gfx-rs/wgpu: A cross-platform, safe, pure-Rust graphics API.

<https://github.com/gfx-rs/wgpu>

3 wgpu: portable graphics library for Rust

<https://wgpu.rs/>

4 5 6 The Benefits of Using Rust for Systems Programming - Sharkbyte

<https://sharkbyte.ca/the-benefits-of-using-rust-for-systems-programming/>

7 In my view, WebGPU is the future of cross-platform graphics. I wrote about this ... | Hacker News

<https://news.ycombinator.com/item?id=29059680>

8 Leverage Rust and wgpu for effective cross-platform graphics - LogRocket Blog

<https://blog.logrocket.com/rust-wgpu-cross-platform-graphics/>

9 10 Potential-Field Source-Surface Models - NSO - National Solar Observatory

<https://nso.edu/data/nisp-data/pfss/>

11 12 High Dynamic Range Rendering | Learn Wgpu

<https://sotrh.github.io/learn-wgpu/intermediate/tutorial13-hdr/>

13 egui\_dock - Rust

[https://docs.rs/egui\\_dock/latest/egui\\_dock/](https://docs.rs/egui_dock/latest/egui_dock/)

14 dear-imgui-wgpu - crates.io: Rust Package Registry

<https://crates.io/crates/dear-imgui-wgpu>

15 Profiling wgpu app with Nsight Graphics. #7761 - GitHub

<https://github.com/gfx-rs/wgpu/discussions/7761>